

Bags, Vectors & Transformers

Foundations of Word Embeddings · Day 2, Part 1

Denise J. Roth, PhD Candidate

Strategic Communication Group · Wageningen University & Research

Welcome Back — Where We Left Off

Yesterday we built the **classic toolkit** for turning text into data. A quick recap:

What we covered

- Text → tokens → clean tokens
- The **document-term matrix**
- TF-IDF weighting
- Dictionary, supervised ML, topic models

And where it broke

- Word order lost
- Synonyms look unrelated
- Sparse, high-dimensional
- No sense of context or meaning

Today we address that last column — starting with the idea of representing meaning itself.

Any questions from yesterday? 

Anything from the notebooks, the document-term matrix, the three approaches, or the limitations we saw?

Now is a good moment — today builds directly on these ideas.

The Problem We Want to Solve

Bag-of-words treats every word as an **isolated symbol**. “happy”, “joyful”, and “concrete” are just three different columns — equally unrelated to one another.

What we *want* is a representation where a word’s **meaning** is captured, so that words with similar meanings have similar representations. The question is:

**Where could “meaning” possibly come from,
if all we have is raw text?**

The answer is one of the most influential ideas in language science.

The Distributional Hypothesis

“You shall know a word by the company it keeps.”

— J.R. Firth, 1957

The **distributional hypothesis** is the idea that words appearing in **similar contexts** tend to have **similar meanings**. Meaning is not an intrinsic label — it can be inferred from **patterns of use**.

This is a genuinely powerful move: it turns an abstract, philosophical question (“what is meaning?”) into something **measurable** from text. If we can measure the contexts a word appears in, we can estimate its meaning.

An Intuition: Guess the Word

Suppose you keep seeing an unfamiliar word — let's call it “**wug**” — in sentences like these:

“I poured myself a cold glass of **wug**.”

“She drinks **wug** with breakfast every morning.”

“The carton of **wug** in the fridge has gone sour.”

You have never been told what “wug” means — but you already have a strong guess. Something like **milk**. How? You inferred it **entirely from context**: the words that surround it (glass, drinks, breakfast, carton, fridge, sour).

That is the distributional hypothesis at work — and it is exactly what embedding methods automate over billions of sentences.

From Contexts to Vectors

If meaning comes from context, we can represent each word by **the company it keeps** — as a **vector of numbers**. Words used in similar contexts get similar vectors.

Bag-of-words vector

- One dimension **per vocabulary word**
- Tens of thousands of dimensions
- Mostly zeros (sparse)
- Dimensions are **just word identities**

Embedding vector

- A few **hundred** dimensions
- **Dense** (mostly non-zero)
- Learned from context
- Dimensions capture **latent features** of meaning

*We call these dense vectors **word embeddings**: each word is “embedded” as a point in a shared meaning-space.*

Words as Points in Space

The key payoff: in embedding space, **distance means something**. Words with similar meanings sit **close together**; unrelated words sit far apart.



“happy”, “joyful”, “cheerful” cluster together; “dog”, “cat”, “rabbit” form a separate cluster. The model was never *told* these groupings — it discovered them from context. (*A 2-D sketch; real embeddings have hundreds of dimensions.*)

A Famous Consequence: Analogies

Because meaning is encoded geometrically, **directions** in the space can carry meaning too. The classic example:

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

The vector that takes you from “man” to “woman” is roughly the *same* vector that takes you from “king” to “queen”. A “gender direction” emerges in the space, along with directions for tense, plurality, even country-to-capital.

This was the result that made embeddings famous around 2013 — and it falls out of the distributional hypothesis, no hand-coding required.

Two Flavors of Embeddings

Static embeddings

- One fixed vector **per word**
- Word2Vec, GloVe
- “bank” has **one** vector, no matter the sentence
- Simple, fast, interpretable

We start here.

Contextual embeddings

- A vector **per word-in-context**
- BERT and friends
- “bank” gets **different** vectors in different sentences
- Powerful, but heavier

We build to this later.

Both rest on the same foundation: meaning from context. The difference is whether context is baked in once, or computed fresh each time.

Foundations: The Takeaways

- ✓ Bag-of-words fails because it treats words as **isolated symbols**
- ✓ The **distributional hypothesis**: a word's meaning can be inferred from the **contexts it appears in**
- ✓ **Embeddings** represent each word as a **dense vector** in a shared meaning-space
- ✓ In that space, **distance** reflects similarity and **directions** can carry meaning (analogies)
- ✓ Two flavors: **static** (one vector per word) and **contextual** (one vector per use)

Next: how do we actually *learn* these vectors from data? We open up Word2Vec and GloVe — and then get our hands on real embeddings in Python.